

Lecture 11: Regression models

FIE463: Numerical Methods in Macroeconomics and Finance using Python

Richard Foltyn
NHH Norwegian School of Economics

April 7, 2026

See GitHub repository for notebooks and data:

<https://github.com/richardfoltyn/FIE463-V26>

Contents

1	Regression models with Scikit-learn	1
1.1	Motivation	1
1.2	Ridge regression	2
1.3	Lasso	10
1.4	Elastic Net	18
1.5	List of functions used in this lecture	22

1 Regression models with Scikit-learn

1.1 Motivation

In the previous lecture, we studied linear models of the form

$$y_i = \mu + \mathbf{x}_i' \boldsymbol{\beta} + \epsilon_i$$

where μ is the intercept (or bias in machine learning lingo), $\mathbf{x}$ is a vector of explanatory variables, and ϵ_i is an additive error term that captures any residual variance in y_i not accounted for by the explanatory variables.

We estimated the coefficients using ordinary least squares (OLS), which is equivalent to minimizing the sum of squared errors,

$$L(\mu, \boldsymbol{\beta}) = \underbrace{\sum_{i=1}^N (y_i - \mu - \mathbf{x}_i' \boldsymbol{\beta})^2}_{\text{Sum of squared errors}} \quad (1)$$

The estimates $(\hat{\mu}, \hat{\boldsymbol{\beta}})$ are then the values that minimize $L(\mu, \boldsymbol{\beta})$,

$$(\hat{\mu}, \hat{\boldsymbol{\beta}}) = \arg \min_{\mu, \boldsymbol{\beta}} L(\mu, \boldsymbol{\beta})$$

While OLS is the best linear unbiased estimator (BLUE) under the usual assumptions about the errors ϵ_i , the estimated $(\hat{\mu}, \hat{\boldsymbol{\beta}})$ may not be the best model parameters when it comes to minimizing *out-of-sample* prediction errors. Whenever our focus is on *prediction* rather than on *inference* about population parameters, we can often do better than using the loss function from (1). To this end, we augment the loss function with penalty terms which prevent overfitting and improve out-of-sample prediction, thus making these models more suitable for machine learning applications.

1.2 Ridge regression

The first model we study is the Ridge regression, which estimates a linear model but adds a penalty term to the loss function:

$$L(\mu, \beta) = \underbrace{\sum_{i=1}^N (y_i - \mu - \mathbf{x}_i' \beta)^2}_{\text{Sum of squared errors}} + \underbrace{\alpha \sum_{k=1}^K \beta_k^2}_{\text{L2 penalty}}$$

Compared to ordinary least squares (OLS) we discussed initially, the penalty term increases the loss, in particular if the estimated coefficients β are large. This term is called an L2 penalty because it corresponds to the (squared) L2 vector norm. In many textbooks, you will see the penalty term written as $\alpha \|\beta\|_2^2$, which is equivalent to the formula used above.

For any $\alpha > 0$, the loss function is increasing in the (absolute) coefficient values, thus the minimum L might be one where the elements of β are smaller than what they would be with OLS. We therefore say that the Ridge regression applies *regularization* or *shrinkage*. Note, however, that the intercept μ is not included in the penalty term, and thus no regularization is applied to it.

Clearly, the regularization strength depends on the value of α . For tiny (or zero) α , the estimated $\hat{\beta}$ will be close (or identical) to OLS, while for large α the estimated coefficients will be compressed towards zero. In this setting, α is a hyperparameter and we can accordingly use cross-validation to find an “optimal” value.

Why would we ever want to use Ridge regression given that OLS is the best linear unbiased estimator? It turns out that regularization can help in scenarios where we have a large number of (potentially multicollinear) regressors, in which case OLS is prone to overfitting.

1.2.1 Example: Polynomial approximation

We illustrate such problems and the benefits of Ridge regression using a highly stylized example. Imagine that we want to approximate a (non-linear) trigonometric function using a high-order polynomial, a setting which is notoriously susceptible to overfitting (also see the previous lecture and workshop for illustrations). Assume that our model is given by

$$y_i = \cos\left(\frac{3\pi}{200}x_i\right) + \epsilon_i$$
$$\epsilon_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 0.5^2)$$

where ϵ_i is an additive, normally-distributed measurement error term with mean 0 and variance $\frac{1}{4}$. For this exercise, we assume that x is uniformly distributed on the interval $[0, 100]$.

The true values of y (without measurement error) are computed using the function `compute_true_y()` which returns y for a given x .

```
[1]: import numpy as np

def compute_true_y(x):
    """
    True trigonometric function (without errors)
    """
    return np.cos(3 * np.pi / 200.0 * x)
```

The following code creates a demo sample with $N = 200$ observations.

```
[2]: from numpy.random import default_rng

def create_trig_sample(N=200, sigma=0.5, rng=None):
    """
    Create trigonometric relationship sample data for Ridge and Lasso.
```

```

Parameters
-----
N : int
    Sample size.
sigma : float
    Standard deviation of the normal error term.
rng : np.random.Generator, optional
    Random number generator.
"""

# Initialize random number generator
if rng is None:
    rng = default_rng(seed=1234)

# Randomly draw explanatory variable x uniformly distributed on [0, 100]
x = rng.uniform(0, 100, size=N)

# Draw errors from normal distribution
epsilon = rng.normal(scale=sigma, size=N)

# Compute y, add measurement error
y = compute_true_y(x) + epsilon

return x, y

```

```
[3]: x, y = create_trig_sample()
```

The next graph visualizes the true relationship and the sample points (x_i, y_i) . For this purpose we first define the function `plot_trig_sample()` which creates the baseline plot used throughout the next few sections.

```
[4]: import matplotlib.pyplot as plt


def plot_trig_sample(x, y):
    """
    Plot the trigonometric relationship sample for Ridge and Lasso
    """
    # Sample scatter plot
    fig, ax = plt.subplots(1, 1, figsize=(5.5, 3.5))
    ax.scatter(x, y, s=20, c='none', edgecolor='steelblue', lw=0.75, label='Sample')

    # Plot true relationship
    xvalues = np.linspace(0.0, 100.0, 101)
    y_true = compute_true_y(xvalues)
    ax.plot(xvalues, y_true, c='black', lw=1.0, label='True function')
    ax.set_xlabel('$x$')
    ax.set_ylabel('$y$')
    ax.legend(loc='upper right')

    return ax

```

```
[5]: _ = plot_trig_sample(x, y)
```



1.2.2 Estimating the Ridge and linear regression models

We now estimate a polynomial approximation where we assume that y is a degree- K polynomial in x , i.e.,

$$y_i \approx \mu + \beta_1 x_i + \beta_2 x_i^2 + \dots + \beta_K x_i^K$$

For this example, we choose an unconventionally high $K = 15$ since we anticipate that this leads to problems with OLS.

A few more steps are required before we can run the Ridge regression:

- We need to create the polynomial in x which we do with the `PolynomialFeatures` transformation we already encountered.
- Moreover, regularization methods can be susceptible to scaling issues, so we need to demean and normalize all input variables, i.e., we make sure that each feature has mean 0 and a variance of 1. We can automate this step using the `StandardScaler` transformation.
- Finally, the estimation step is performed by the `Ridge` class.
- We build a `Pipeline` that combines these transformations together with the estimation step. We use the function `make_pipeline()` to simplify this step.

To run the Ridge regression, we also need to specify the regularization strength α which we set to $\alpha = 3$ for illustration.

```
[6]: from sklearn.linear_model import Ridge
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.pipeline import make_pipeline

# Polynomial degree
degree = 15

# Penalty strength to use for Ridge
alpha = 3

# Build pipeline of transformations and Ridge estimation.
# We create the polynomial transformation w/o the intercept so we
# need to include an intercept in the fitting step.
pipe_ridge = make_pipeline(
    PolynomialFeatures(
```

```

        degree=degree,
        include_bias=False
    ),
    StandardScaler(),           # standardize features
    Ridge(alpha=alpha, fit_intercept=True) # Fit Ridge regression
)

# Make sure X is a matrix
X = x[:, None]

# Fit model
pipe_ridge.fit(X, y)

```

```

[6]: Pipeline(steps=[('polynomialfeatures',
                      PolynomialFeatures(degree=15, include_bias=False)),
                     ('standardscaler', StandardScaler()),
                     ('ridge', Ridge(alpha=3))])

```

It is instructive to estimate the same model using linear regression and compare the results:

```

[7]: from sklearn.linear_model import LinearRegression

# Create pipeline for linear model
pipe_lr = make_pipeline(
    PolynomialFeatures(
        degree=degree,
        include_bias=False
    ),
    StandardScaler(),
    LinearRegression(fit_intercept=True)
)

# Make sure X is a matrix
X = x[:, None]

# Fit model
pipe_lr.fit(X, y)

```

```

[7]: Pipeline(steps=[('polynomialfeatures',
                      PolynomialFeatures(degree=15, include_bias=False)),
                     ('standardscaler', StandardScaler()),
                     ('linearregression', LinearRegression())])

```

Note that in theory you don't need a `StandardScaler()` for the OLS regression, as scaling the explanatory variable just scales the estimated coefficient in the other direction. However, this is only true for the linear algebra expression you find in econometrics textbooks. In practice, your regression matrix may become ill-conditioned if the explanatory variables are of vastly different magnitudes, as is the case with our 15th-order polynomial. Applying a `StandardScaler()` helps to mitigate numerical instability in such scenarios, and doesn't do any harm if no such issues are present. However, you have to be careful about the interpretation of the estimated coefficients when scaling is applied.

To illustrate the difference for this artificial example, we add the model predictions from the Ridge and linear regressions to the sample scatter plot we created earlier.

```

[8]: # Grid on which to evaluate predictions
xvalues = np.linspace(np.amin(x), np.amax(x), 100)

# Predicted values from linear regression
y_pred_lr = pipe_lr.predict(xvalues[:, None])

# Predicted values from Ridge regression
y_pred_ridge = pipe_ridge.predict(xvalues[:, None])


```

```
[9]: # Plot sample and true relationship
ax = plot_trig_sample(x, y)

# Linear regression prediction
ax.plot(xvalues, y_pred_lr, c='purple', alpha=0.7, label='Linear regression')

# Ridge prediction
ax.plot(xvalues, y_pred_ridge, c='darkorange', lw=2.0, label='Ridge')

ax.set_title(rf'Ridge regression with  $\alpha = \{alpha:.4g\}$  vs OLS benchmark')
_ = ax.legend()
```



As the graph shows, the OLS model vastly overfits the data, which is what we would expect in this setting. Conversely, the prediction of the Ridge regression is reasonably close to the true function and much better behaved despite the high polynomial degree used here.

Your turn.

Rerun the above estimation for Ridge and OLS, but remove the `StandardScaler()` from the pipeline. What happens to the predicted values?

1.2.3 Ridge path

To gain some intuition for what is going on, it is instructive to plot the estimated Ridge coefficients for a whole range of α values as we do in the code below. This plot is called the *Coefficient path* or *Ridge path* in the context of the Ridge regression.

Note that we choose the grid of α to be uniformly spaced in logs since we want to zoom in on what happens when α is small. This is accomplished by using `np.logspace()` instead of `np.linspace()`.

```
[10]: # Create grid of alphas spaced uniformly in logs on [5e-3, 1000]
alphas = np.logspace(start=np.log10(5.0e-3), stop=np.log10(1000), num=100)

# Re-create pipeline w/o Ridge estimator, estimation step differs for each alpha
transform = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False), StandardScaler()
)
```

```

# Create polynomial features
X_trans = transform.fit_transform(x[:, None])

# Array to store coefficients for all alphas
coefs = np.empty((len(alphas), X_trans.shape[1]))

# Estimate Ridge for each alpha, store fitted coefficients
for i, alpha in enumerate(alphas):
    ridge = Ridge(alpha=alpha, fit_intercept=True)

    # Fit model for given alpha
    ridge.fit(X_trans, y)

    coefs[i] = ridge.coef_

```

The following plot shows each coefficient (one color corresponds to one coefficient) for different values of α on the x-axis. Note that the x-axis is plotted on a $\log_{10}$ scale which allows us to zoom in on smaller values of α .


```

[11]: # Plot coefficient arrays against penalty strength
plt.figure(figsize=(6,4))

plt.plot(alphas, coefs, lw=1.0)

plt.xscale('log', base=10)
plt.axhline(0.0, ls='--', lw=0.75, c='black')
plt.xlabel(r'Regularization strength  $\alpha$  (log scale)')
plt.ylabel('Coefficient value')
plt.title('Ridge coefficients as function of regularization strength')
_ = plt.legend([rf'$\beta_{\{i\}}$' for i in range(degree)], ncols=5, loc='lower right')

```



For small α (on the left) the estimated coefficients are close to the (standardized) OLS coefficients, but their values shrink towards zero as α increases. For a very large $\alpha = 10^3$ the estimated coefficients are basically zero since the penalty dominates the sum of squared errors in the loss function.

1.2.4 Tuning the regularization parameter via cross-validation

In the previous example, we picked an arbitrary regularization strength α when fitting the Ridge regression. In applied work, we would want to tune α (which is a hyperparameter) using cross-validation instead. To this end, we could use the generic cross-validation functionality we studied in the previous lecture since that one works for all types of estimators. However, `scikit-learn` implements a cross-validation class specifically for Ridge regression called `RidgeCV` which we use in the code below to find an optimal α .

```
[12]: from sklearn.linear_model import RidgeCV

# RidgeCV does not support pipelines, so we need to transform x before
# cross-validation.
transform = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False), StandardScaler()
)

# Create standardized polynomial features
X_trans = transform.fit_transform(x[:, None])

# Candidate alphas used for cross-validation, spaced uniformly in logs to get
# denser grid for small alphas.
N_alphas = 100
alphas = np.logspace(start=np.log10(1.0e-5), stop=np.log10(5), num=N_alphas)

# Create and run Ridge cross-validation
rcv = RidgeCV(alphas=alphas, store_cv_results=True).fit(X_trans, y)
```

Note that we don't need to specify the number of folds to be used for cross-validation (this defaults to `cv=None` when creating `RidgeCV`). In this case, `RidgeCV` performs the more efficient leave-one-out cross-validation instead of splitting the sample into `n` folds.

By default, `RidgeCV` uses the (negative) mean squared error (MSE) to find the best α which we can then recover from the `alpha_` attribute.

```
[13]: # Recover best alpha that minimizes MSE
alpha_best = rcv.alpha_

# Best MSE is stored as negative score!
MSE_best = - rcv.best_score_

print(f'Best alpha: {alpha_best:.3g} (MSE: {MSE_best:.3g})')
```

Best alpha: 0.00299 (MSE: 0.25)

Because we fitted `RidgeCV` with the argument `store_cv_results=True`, the fitted object stores the MSE for each α in `cv_results_` which we can use to plot the MSE as a function of α . The dimension of `cv_results_` is `(Nobs, Nalphas)`. Note that this attribute is only available with leave-one-out cross-validation, but not when manually specifying the number of cross-validation folds.

```
[14]: # Compute average MSE for each alpha value across all leave-one-out splits
mse_mean = np.mean(rcv.cv_results_, axis=0)


# Index of MSE-minimizing alpha
imin = np.argmin(mse_mean)
```

```
[15]: import matplotlib.pyplot as plt

# Plot MSE against alphas, highlight minimum MSE
plt.plot(alphas, mse_mean)
plt.xlabel(r'Regularization strength $\alpha$ (log scale)')
plt.ylabel('Cross-validated MSE')
```



```
plt.scatter(alphas[imin], mse_mean[imin], s=15, c='black', zorder=100)
plt.axvline(alphas[imin], ls=':', lw=0.75, c='black')
plt.xscale('log')
_ = plt.title('Validation curve for Ridge regression')
```



Now that we have identified the optimal α , we can re-fit the Ridge regression and plot the prediction from this model. Note that this is not strictly necessary as the return value of `RidgeCV` can be used to do prediction based on coefficients estimated for the best α , but because `RidgeCV` does not support pipelines, we'd have to apply any transformations manually before doing so.

```
[16]: # Create pipeline with Ridge estimator using optimal alpha
pipe_ridge = make_pipeline(
    PolynomialFeatures(
        degree=degree,
        include_bias=False
    ),
    StandardScaler(),
    Ridge(alpha=alpha_best, fit_intercept=True)
)

# Fit Ridge regression with optimal alpha
pipe_ridge.fit(x[:, None], y)
```

```
[16]: Pipeline(steps=[('polynomialfeatures',
    PolynomialFeatures(degree=15, include_bias=False)),
    ('standardscaler', StandardScaler()),
    ('ridge', Ridge(alpha=0.0029875157525272996))])
```

```
[17]: # Grid on which to evaluate predictions
xvalues = np.linspace(np.amin(x), np.amax(x), 100)


# Predicted values from Ridge regression
y_pred = pipe_ridge.predict(xvalues[:, None])
```

```
[18]: # Plot sample and true relationship
ax = plot_trig_sample(x, y)

# Plot predicted values from cross-validated Ridge regression
ax.plot(xvalues, y_pred, c='darkorange', lw=2.0, label='Ridge')

ax.set_title(rf'Ridge regression with optimal  $\alpha$  = {alpha_best:.4g}')
```

```
_ = ax.legend()
```



Your turn.

Rerun the whole Ridge example with a smaller sample size of $N = 50$. What happens to the optimal cross-validated penalty parameter α ?

1.3 Lasso

Another widely used estimator with shrinkage is LASSO (least absolute shrinkage and selection operator) which adds an L1 penalty term to the OLS objective function:

$$L(\mu, \beta) = \underbrace{\sum_{i=1}^N (y_i - \mu - \mathbf{x}_i' \beta)^2}_{\text{Sum of squared errors}} + \alpha \underbrace{\sum_{k=1}^K |\beta_k|}_{\text{L1 penalty}}$$

As with the Ridge regression, this additional term penalizes large coefficient values. This term is called an L1 penalty because it corresponds to the L1 vector norm which can equivalently be written as $\alpha \|\beta\|_1$.

While the objective looks very similar to Ridge regression, using the L1 instead of the L2 norm can produce much more parsimonious models because many coefficients end up being exactly zero and the corresponding features are thus eliminated from the model. We will see this in the example below.

1.3.1 Estimating the Lasso and linear regression models

We apply Lasso to the same random sample as in the section on Ridge which allows us to compare the two methods. The following code recreates that data, making the same functional form and distributional assumptions as in the previous section.

```
[19]: x, y = create_trig_sample()
```

As with Ridge, we need to standardize the explanatory variables before fitting Lasso. The code below repeats these steps, but we now use `Lasso` to perform the model estimation. For now, we set the regularization strength to $\alpha = 0.0075$ and will use cross-validation to find the optimal value later.

Note that for Lasso it might be necessary to increase the number of iterations by increasing the `max_iter` parameter (from the default of 1,000).

```
[20]: from sklearn.linear_model import Lasso
      from sklearn.preprocessing import PolynomialFeatures, StandardScaler
      from sklearn.pipeline import make_pipeline

      # Polynomial degree
      degree = 15

      # Penalty strength to use for Lasso
      alpha = 0.0075

      # Build pipeline of transformations and Lasso estimation.
      # We create the polynomial transformation w/o the intercept so we
      # need to include an intercept in the fitting step.
      pipe_lasso = make_pipeline(
          PolynomialFeatures(
              degree=degree,
              include_bias=False
          ),
          StandardScaler(),
          Lasso(alpha=alpha, fit_intercept=True, max_iter=10000)
      )

      # Make sure X is a matrix
      X = x[:, None]

      pipe_lasso.fit(X, y)
```

```
[20]: Pipeline(steps=[('polynomialfeatures',
                        PolynomialFeatures(degree=15, include_bias=False)),
                        ('standardscaler', StandardScaler()),
                        ('lasso', Lasso(alpha=0.0075, max_iter=10000))])
```

For completeness, let's also recreate the linear regression estimation and plot the model predictions alongside the Lasso.

```
[21]: from sklearn.linear_model import LinearRegression

      # Create pipeline for linear model
      pipe_lr = make_pipeline(
          PolynomialFeatures(
              degree=degree,
              include_bias=False
          ),
          StandardScaler(),
          LinearRegression(fit_intercept=True)
      )

      # Make sure X is a matrix
      X = x[:, None]

      pipe_lr.fit(X, y)
```

```
[21]: Pipeline(steps=[('polynomialfeatures',
                        PolynomialFeatures(degree=15, include_bias=False)),
                        ('standardscaler', StandardScaler()),
                        ('linearregression', LinearRegression())])
```

The following code plots the sample data, the true functional relationship, and the predictions from the linear regression and Lasso models.

```
[22]: # Grid on which to evaluate predictions
xvalues = np.linspace(np.amin(x), np.amax(x), 100)

# Predicted values from linear regression
y_pred_lr = pipe_lr.predict(xvalues[:, None])


# Predicted values from Lasso regression
y_pred_lasso = pipe_lasso.predict(xvalues[:, None])

[23]: ax = plot_trig_sample(x, y)

# Linear regression prediction
ax.plot(xvalues, y_pred_lr, c='purple', alpha=0.7, label='Linear regression')

# Lasso prediction
ax.plot(xvalues, y_pred_lasso, c='darkorange', lw=2.0, label='Lasso')

ax.set_title(rf'Lasso with $\alpha$ = {alpha:.4g} vs OLS benchmark')
_ = ax.legend()
```



You might be wondering why we chose $\alpha = 0.0075$ whereas we initially used $\alpha = 3$ for the Ridge regression. The reason is that the scikit-learn implementation of Lasso uses a slightly different loss function than the one given above, namely

$$L(\mu, \beta) = \frac{1}{2N} \sum_{i=1}^N (y_i - \mu - \mathbf{x}_i' \beta)^2 + \alpha \|\beta\|_1$$

which scales the sum of squared errors term by a factor of $(2N)^{-1}$. This makes no difference for the optimization, but changes the interpretation of the regularization strength α compared to Ridge regression. For our sample size of $N = 200$, an α_{Ridge} used for the Ridge estimator approximately corresponds to $\alpha_{\text{Lasso}} = \frac{\alpha_{\text{Ridge}}}{400}$ when plugged into Lasso.

1.3.2 Lasso path

As for the Ridge regression, we can visualize how the estimated Lasso coefficients change as we vary the penalty parameter. This plot is called the *Lasso path*.

To compute the path, we fit the model for different values of α on the interval $[10^{-3}, 1]$ where we again space the α uniformly in logs.

```
[24]: # Create grid of alphas spaced uniformly in logs on [1e-3, 1]
alphas = np.logspace(start=np.log10(1.0e-3), stop=np.log10(1.0), num=100)

# Re-create pipeline w/o Lasso estimator, estimation step differs for each alpha
transform = make_pipeline(
    PolynomialFeatures(
        degree=degree,
        include_bias=False
    ),
    StandardScaler()
)

# Create polynomial features
X_trans = transform.fit_transform(x[:, None])

# Array to store coefficients for all alphas
coefs = np.empty((len(alphas), X_trans.shape[1]))


# Estimate Lasso for each alpha, store fitted coefficients
for i, alpha in enumerate(alphas):
    lasso = Lasso(alpha=alpha,
        fit_intercept=True,
        max_iter=100_000
    )

    # Fit model for given alpha
    lasso.fit(X_trans, y)

    coefs[i] = lasso.coef_
```

Plotting the fitted coefficients against α on a log scale looks as follows:


```
[25]: plt.figure(figsize=(6, 4))
plt.plot(alphas, coefs, lw=1.0)
plt.xscale('log', base=10)
plt.axhline(0.0, ls='--', lw=0.75, c='black')
plt.xlabel(r'Regularization strength  $\alpha$  (log scale)')
plt.ylabel('Coefficient value')
plt.title('Lasso coefficients as function of regularization strength')
_ = plt.legend([rf' $\beta_{\{i\}}$ ' for i in range(degree)], ncols=5)
```



The graph shows that the coefficient estimates are quite different from what we obtained from the Ridge estimator. In fact, most of them are exactly zero for most values of α . We highlight this result in the graph below which plots the number of nonzero coefficients against α .

```
[26]: # Number of nonzero coefficients for each alpha.
nonzero = np.sum(np.abs(coefs) > 1.0e-6, axis=1).astype(int)
```

```
[27]: # Plot number of nonzero coefficients against alpha
plt.plot(alphas, nonzero, lw=1.5, c='steelblue')
plt.xscale('log', base=10)
plt.yticks(np.arange(0, np.amax(nonzero) + 1))
plt.xlabel(r'Regularization strength  $\alpha$  (log scale)')
_ = plt.title('Number of nonzero coefficients')
```



Clearly, in this case the model estimated by Lasso is substantially less complex than the linear regression

or even Ridge regression. For most values of α , only 2–3 features out of the original 15 are retained in the model.

Note that it is possible to compute the path of Lasso coefficients along the grid of α using the convenience function `lasso_path()`. This function does not support pipelines, so we have to apply any transformations and standardizations to X ourselves:


```
[28]: from sklearn.linear_model import lasso_path

# Re-create pipeline w/o Lasso estimator to be used with lasso_path()
transform = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False), StandardScaler()
)

# Create polynomial features
X_trans = transform.fit_transform(x[:, None])

# Compute Lasso path automatically
alphas, coefs, _ = lasso_path(X_trans, y, alphas=alphas, max_iter=100_000)
```

```
[29]: plt.figure(figsize=(6, 4))
plt.plot(alphas, coefs.T, lw=1.0)
plt.xscale('log', base=10)
plt.axhline(0.0, ls='--', lw=0.75, c='black')
plt.xlabel(r'Regularization strength  $\alpha$  (log scale)')
plt.ylabel('Coefficient value')
plt.title('Lasso coefficients as function of regularization strength')
_ = plt.legend([rf'$\beta_{\{i\}}$' for i in range(degree)], ncol=5)
```



As you can see, the estimated coefficients along the path are not exactly the same as when doing it manually, which might be because the two approaches call the underlying coordinate descent optimizer in different ways.

1.3.3 Tuning the regularization parameter via cross-validation

In the previous example, we picked an arbitrary regularization strength α when fitting the Lasso. In this section, we again find an optimal α using cross-validation. Just like in the case of Ridge regression, `scikit-learn` implements a cross-validation class specifically for Lasso called `LassoCV` which we use in the code below to find an optimal α .

`LassoCV` optionally accepts a grid of candidate α values just like `RidgeCV`, but the default way to run cross-validation is to specify the fraction $\epsilon = \frac{\alpha_{\min}}{\alpha_{\max}}$ (default: 10^{-3}) and the grid size (default: 100). `LassoCV` then automatically determines an appropriate grid which is stored in the `alphas_` attribute once fitting is complete. Moreover, we use the `cv` argument to set the desired number of CV folds (default: 5).

```
[30]: from sklearn.linear_model import LassoCV

# LassoCV does not support pipelines, so we need to transform x before
# cross-validation.
transform = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False), StandardScaler()
)

# Create standardized polynomial features
X_trans = transform.fit_transform(x[:, None])

# Create and run Lasso cross-validation, use defaults for eps and n_alphas
lcv = LassoCV(max_iter=100_000, cv=10).fit(X_trans, y)
```

After fitting, we can recover the best alpha from the `alpha_` attribute and the MSE for each α on the grid and each CV fold from the `mse_path_` attribute.

```
[31]: # Recover best alpha that minimizes MSE
alpha_best = lcv.alpha_

# MSE for each alpha, averaged over folds
mse_mean = np.mean(lcv.mse_path_, axis=1)

# Index of min. MSE
imin = np.argmin(mse_mean)

# Store best MSE
mse_best = mse_mean[imin]

print(f'Best alpha: {alpha_best:.4g} (MSE: {mse_best:.3g})')
```

Best alpha: 0.0007855 (MSE: 0.251)

After fitting, the `LassoCV` instance stores the estimated parameters at the optimal α in the attribute `coef_`. We can use this vector to compute the number of features (explanatory variables) with nonzero coefficients:

```
[32]: n_features = np.sum(np.abs(lcv.coef_) > 0)
print(f'Number of nonzero coefficients: {n_features}')
```

Number of nonzero coefficients: 6

The next plot visualizes the average MSE over the entire range of candidate α values which are stored in the `alphas_` attribute of the `LassoCV` instance:


```
[33]: import matplotlib.pyplot as plt

# Recover grid of alphas used for CV
alphas = lcv.alphas_

# Plot MSE against alphas, highlight minimum MSE
```



```
plt.plot(alphas, mse_mean)
plt.xlabel(r'Regularization strength  $\alpha$  (log scale)')
plt.ylabel('Cross-validated MSE')
plt.scatter(alphas[imin], mse_mean[imin], s=15, c='black', zorder=100)
plt.axvline(alphas[imin], ls=':', lw=0.75, c='black')
plt.xscale('log')
_ = plt.title('Validation curve for Lasso')
```



Now that we have identified the optimal α , we can re-fit the Lasso and plot the prediction from this model.

```
[34]: # Create pipeline with Lasso using optimal alpha
pipe_lasso = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False),
    StandardScaler(),
    Lasso(alpha=alpha_best, fit_intercept=True, max_iter=100_000),
)

pipe_lasso.fit(x[:, None], y)
```

```
[34]: Pipeline(steps=[('polynomialfeatures',
    PolynomialFeatures(degree=15, include_bias=False)),
    ('standardscaler', StandardScaler()),
    ('lasso',
    Lasso(alpha=np.float64(0.0007855307824807166),
    max_iter=100000))])
```

Finally, we visually compare the true model to the Lasso prediction using the optimal value of α .


```
[35]: # Grid on which to evaluate predictions
xvalues = np.linspace(np.amin(x), np.amax(x), 100)

# Predicted values from Lasso regression
y_pred = pipe_lasso.predict(xvalues[:, None])
```

```
[36]: # Plot sample and true relationship
ax = plot_trig_sample(x, y)

# Plot prediction from optimal Lasso model
ax.plot(xvalues, y_pred, c='darkorange', lw=2.0, label=r'Lasso')
```

```
ax.set_title(rf'Lasso with optimal $\alpha$ = {alpha_best:.4g}')
_ = ax.legend()
```



1.4 Elastic Net

Elastic Net is a regularized linear regression model that combines both the L1 and L2 penalties of the Lasso and Ridge methods. The objective function is given by:

$$L(\mu, \beta) = \underbrace{\sum_{i=1}^N (y_i - \mu - \mathbf{x}_i' \beta)^2}_{\text{Sum of squared errors}} + \underbrace{\alpha \rho \sum_{k=1}^K |\beta_k|}_{\text{L1 penalty}} + \underbrace{\alpha(1 - \rho) \sum_{k=1}^K \beta_k^2}_{\text{L2 penalty}}$$

Note: As with Lasso, the `scikit-learn` implementation uses a slightly different loss function scaled by $\frac{1}{2N}$ and applies a $\frac{1}{2}$ factor to the L2 penalty, but the qualitative trade-off remains the same.

By incorporating both penalties, Elastic Net allows for learning a sparse model where few of the weights are nonzero (like Lasso), while still maintaining the regularization properties of Ridge regression. The hyperparameter $\rho \in [0, 1]$ (which is called `l1_ratio` in `scikit-learn`) controls the convex combination of L1 and L2 penalties: - When $\rho = 1$, the model is equivalent to Lasso. - When $\rho = 0$, the model is equivalent to Ridge regression.

1.4.1 Estimating the Elastic Net and linear regression models

We use the same simulated dataset from the previous sections to fit the Elastic Net model.

```
[37]: x, y = create_trig_sample()
```

We first illustrate the estimation for a fixed set of hyperparameters. Here, we specify `alpha = 0.01` and `l1_ratio = 0.5` (an equal mix of L1 and L2 penalties). We also compute the ordinary least squares (OLS) benchmark using the `LinearRegression` class for comparison.

```
[38]: from sklearn.linear_model import ElasticNet, LinearRegression
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.pipeline import make_pipeline
```

```

# Polynomial degree
degree = 15

# Build pipeline for Elastic Net
pipe_enet = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False),
    StandardScaler(),
    ElasticNet(alpha=0.01, l1_ratio=0.5, fit_intercept=True, max_iter=100_000)
)

# Build pipeline for OLS (Linear Regression) benchmark
pipe_ols = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False),
    StandardScaler(),
    LinearRegression(fit_intercept=True)
)

# Make sure X is a matrix
X = x[:, None]

# Fit both models
pipe_enet.fit(X, y)
pipe_ols.fit(X, y)

```

```

[38]: Pipeline(steps=[('polynomialfeatures',
                        PolynomialFeatures(degree=15, include_bias=False)),
                      ('standardscaler', StandardScaler()),
                      ('linearregression', LinearRegression())])

```

We now compare the predicted values from both models against the true underlying function.

```

[39]: import numpy as np

# Grid on which to evaluate predictions
xvalues = np.linspace(np.amin(x), np.amax(x), 100)


# Predicted values
y_pred_enet = pipe_enet.predict(xvalues[:, None])
y_pred_ols = pipe_ols.predict(xvalues[:, None])

# Plot sample and true relationship
ax = plot_trig_sample(x, y)

# Plot predictions
ax.plot(xvalues, y_pred_ols, c='purple', lw=1.5, ls='-', label='OLS')
ax.plot(xvalues, y_pred_enet, c='darkorange', lw=2.0, label='Elastic Net')

ax.set_title(r'Elastic Net vs. OLS benchmark')
_ = ax.legend()

```



1.4.2 Tuning the regularization parameters via cross-validation

Unlike Ridge and Lasso, Elastic Net has two hyperparameters to tune: α and the L1 ratio (ρ). `scikit-learn` provides the `ElasticNetCV` class, which can perform a grid search over both parameters simultaneously.

We define an array of possible values for the L1 ratio and let `ElasticNetCV` find the optimal combination.

```
[40]: from sklearn.linear_model import ElasticNetCV

# Create pipeline for transformations
transform = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False),
    StandardScaler()
)
X_trans = transform.fit_transform(x[:, None])

# Candidate values for L1 ratio
l1_ratios = [0.01, 0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1.0]

# Candidate alphas
alphas = np.logspace(start=np.log10(1.0e-5), stop=np.log10(5), num=100)

# Create and run Elastic Net cross-validation
enet_cv = ElasticNetCV(
    l1_ratio=l1_ratios,      # grid of candidate L1 ratios
    alphas=alphas,          # grid of candidate alphas
    cv=5,                   # use 5-fold cross-validation
    max_iter=10_000_000,    # increase max iterations to ensure convergence
    n_jobs=-1               # use all available CPU cores for parallel processing
)

# Run cross-validation to find optimal alpha and L1 ratio
enet_cv.fit(X_trans, y)

# Recover optimal hyperparameters
alpha_best = enet_cv.alpha_
```

```

l1_ratio_best = enet_cv.l1_ratio_

print(f'Best alpha: {alpha_best:.3g}')
print(f'Best L1 ratio: {l1_ratio_best:.3g}')

```

```

/home/richard/.conda/envs/FIE463/lib/python3.13/site-
packages/sklearn/linear_model/_coordinate_descent.py:701: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.744e-02, tolerance: 1.179e-02
  model = cd_fast.enet_coordinate_descent_gram(

Best alpha: 0.00262
Best L1 ratio: 0.01

```

As you can see from the output, the chosen `l1_ratio` of 0.01 is quite small, meaning that the Elastic Net puts much more weight on the Ridge penalty term than on the Lasso one.

Note that you might have to massively increase the number of maximum iterations for the algorithm to converge since the default is only 1,000. However, with more iterations the cross-validation can take considerable amounts of time. You can speed up the process by telling `scikit-learn` to use more than one core. For example, in the above code we specify `n_jobs=-1`, which tells the cross-validation to use *all* available cores on your machine. Note that using too many cores might exhaust the memory (RAM) available on your computer.

With the optimal hyperparameters identified, we can reconstruct the pipeline and visualize the final, optimized Elastic Net prediction.

```

[41]: # Create pipeline with optimized Elastic Net
pipe_enet_opt = make_pipeline(
    PolynomialFeatures(degree=degree, include_bias=False),
    StandardScaler(),
    ElasticNet(
        alpha=alpha_best, l1_ratio=l1_ratio_best, fit_intercept=True, max_iter=100_000
    ),
)


# Fit optimized model
pipe_enet_opt.fit(x[:, None], y)
y_pred_opt = pipe_enet_opt.predict(xvalues[:, None])

# Plot sample and true relationship
ax = plot_trig_sample(x, y)

# Plot optimal prediction
ax.plot(xvalues, y_pred_opt, c='darkorange', lw=2.0, label='Elastic Net')

ax.set_title(
    rf'Elastic Net with optimal $\alpha$={alpha_best:.4g}, $l_1$-ratio$={l1_ratio_best:.3g}'
)
_ = ax.legend()

```



1.5 List of functions used in this lecture

The following list contains the central functions covered in this lecture. Each function name is linked to the official API documentation which you can consult for more details regarding function arguments and return values. The [go to section] links to the relevant section in this notebook.

1.5.1 NumPy array functions

- `logspace()`: Function for creating numbers spaced evenly on a log scale. [\[go to section\]](#)

1.5.2 Scikit-learn linear models

- `Ridge()`: Linear least squares with L2 regularization. [\[go to section\]](#)
- `RidgeCV()`: Ridge regression with built-in cross-validation. [\[go to section\]](#)
- `Lasso()`: Linear model with L1 regularization (shrinkage). [\[go to section\]](#)
- `lasso_path()`: Function for computing the Lasso path with coordinate descent. [\[go to section\]](#)
- `LassoCV()`: Lasso linear model with iterative fitting along a regularization path. [\[go to section\]](#)
- `ElasticNet()`: Linear regression with combined L1 and L2 priors as regularizer. [\[go to section\]](#)
- `ElasticNetCV()`: Elastic Net model with iterative fitting along a regularization path. [\[go to section\]](#)